

Dialogue Benchmark Generation from Knowledge Graphs with Cost-Effective Retrieval-Augmented LLMs

Reham Omar, Omij Mangukiya, Essam Mansour

SIGMOD 2025


November 9, 2025

Outline

- 1 Background
- 2 Challenges & Contributions
- 3 Method / Architecture
- 4 Dialogue Context Extraction
- 5 Multi-stage Dialogue Generation
- 6 Conclusion & Limitations

1.1 What is a Dialogue Benchmark?

- A dataset of multi-turn dialogues used to evaluate dialogue systems.
- Often represented as sequences (Q, A) or full dialogues with context and gold answers.
- **Applications:** educational chatbots, domain-specific assistants, benchmarking research models.



1.2 Limitations of Prior Methods

Traditional (document-based) approaches

- Manual authoring (CoQA, QuAC): high cost, low scalability.
- Template-based systems: brittle and require per-KG templates.

KG-based approaches

- Rule/template systems (CSQA, Head-to-Tail, Maestro): heavy preprocessing, limited dialogue support.
- Do not handle hallucinations or provide end-to-end dialogue generation with validations.


2.1 Why LLM + RAG? Challenges

- **Efficient KG retrieval:** scale — millions/billions of entities.
- **Prompt design:** complex prompts can overwhelm LLMs; few-shot vs zero-shot tradeoffs.
- **Hallucination and synthesis:** LLMs may invent facts not in the KG.
- **Cross-model consistency:** want approach working with GPT-4o, Gemini, Llama-3, etc.

2.2 Contributions / Chatty-Gen Overview

- A fully automated, multi-stage RAG pipeline (Chatty-Gen) for KG-grounded dialogue benchmark generation.
- Key features: **type-aware sampling, textual entity labels, summarized subgraphs, multi-stage generation, assertion validation.**
- Addresses challenges:
 - **retrieval efficiency** (type-aware sampling)
 - **prompt complexity** (summarized subgraphs & multi-stage generation)
 - **hallucination** (assertion-based validators)
 - **cross-LLM compatibility**

3 Chatty-Gen Architecture



(Use Figure 2 from paper: architecture diagram)

- Phase A: **Dialogue Context Extraction** (node type selection, textual representation, seed sampling, subgraph extraction)
- Phase B: **Dialogue Generation** (subgraph summarization, question generation, SPARQL generation, dialogue generation)

4.1 Algorithm 1: Representative Node Type Selection

Goal: Efficiently select representative entity types and determine sampling size per type.

Key Ideas:

- Reduce cost by operating on **types**, not individual entities.
- Remove:
 - Metadata types (e.g., ontology or administrative nodes)
 - Low-frequency types (threshold R)
 - Shadowed parent types (threshold S)
- Output a pruned and meaningful type list with sample sizes.

Algorithm 1 Node Type Selection

Input: endpoint : SPARQL endpoint, m : number of Dialogues, domain : prefixes of KG, R : Rare types threshold, S : Shadowed parents threshold

Output: dist : A map of node type to the number of entities

```
1:  $\text{dist}, \text{type\_ratio} \leftarrow \{\}$ 
2:  $\text{types}, \text{count} \leftarrow \text{getKGNodeTypes}(\text{endpoint}, \text{domain})$ 
3:  $\text{total} \leftarrow \text{Sum}(\text{count})$ 
4: for every  $\langle t, c \rangle \in \langle \text{types}, \text{count} \rangle$  do
5:    $\text{type\_ratio}[t] \leftarrow c/\text{total}$ 
6: end for
7:  $\text{type\_ratio} \leftarrow \text{removeRareTypes}(\text{type\_ratio}, R)$ 
8:  $\text{type\_ratio} \leftarrow \text{removeShadowedTypes}(\text{type\_ratio}, S)$ 
9: for every  $\langle t, \text{ratio} \rangle \in \text{type\_ratio}$  do
10:    $\text{dist}[t] \leftarrow \text{ratio} * m$ 
11: end for
12: Return  $\text{dist}$ 
```

Algorithm 1: Node Type Selection

4.1 Algorithm 1: Steps

① Query types & counts.

- Run a SPARQL aggregation on `rdf:type` to get (type, count).

② Filter metadata / domain.

- Remove ontology/meta types and restrict to desired namespace prefixes.

③ Remove rare types (threshold R).

- Compute $\text{ratio}(t) = \frac{\text{count}(t)}{\sum_{t'} \text{count}(t')}$.
- Drop types with $\text{ratio}(t) < R$ (e.g. $R = 1\%$).

④ Remove shadowed parent types (threshold S).

- For parent p and child c , compute $\text{cover}(p, c) = \frac{\text{bothCount}(p, c)}{\text{parentCount}(p)}$.
- If $\text{cover} > S$ (e.g. $S = 99\%$), drop the parent.

⑤ Allocate sample budget.

- Given total target dialogues m : $\text{dist}[t] = \text{round}(\text{ratio}(t) \times m)$.
- Resolve leftover by assigning to top-ranked types.

4.1 Algorithm 1: Complexity & Example

Complexity: $O(q_{\text{cost}} + \tau)$ — q_{cost} for SPARQL; τ = number of types.

Tiny example:

- **Type removal (rare):** Person=10,000; Place=2,000; Reservoir=2.
Total=12,002. With $R = 1\%$ remove Reservoir. For $m = 120$: Person $\rightarrow 100$, Place $\rightarrow 20$.
- **Shadowed parent example:**
 - parent = Creator, parentCount = 1,000
 - child = Person, childCount = 1200
 - bothCount (entities labeled both Creator & Person) = 995
 - $\text{cover}(\text{Creator}, \text{Person}) = 995 / 1000 = 0.995 = 99.5\% > S = 99\% \Rightarrow \text{drop parent 'Creator'}$.

4.2 Textual Entity Representation: Problem & Idea

Problem.

- KG entities are URIs (e.g. .../Q123456), not readable names.
- URI local-names may be IDs or include disambiguation tokens like “(film)” — unreliable as natural labels.

Traditional fix (limits).

- Rule-based mapping (use dbp:title or rdfs:label) fails across heterogeneous KGs (some use skos:prefLabel, foaf:name, etc.).

Key idea.

- Use an LLM to *select the best textual predicate* for each node type from candidate literal predicates — then use its value as the entity name in prompts.

4.2 Textual Entity Representation: Method

Method (per type/entity).

- ① Retrieve candidate literal predicates for a sample entity of the type (predicates with literal objects).
- ② Prompt the LLM: ask which predicate best represents the entity's human-readable name.
- ③ Save mapping $\text{type} \mapsto$ chosen predicate; when generating, extract that predicate's literal as the display name (replace URI).

Benefits. Portable across KGs, improves readability, reduces hallucination by giving LLM natural entity names.

4.2 Textual Entity Representation: *Type*-Level Selection

SPARQL (list literal predicates for a *type*, e.g., dbo:Film):

SPARQL (type-level candidates)

```
SELECT DISTINCT ?p WHERE {  
    ?s a dbo:Film; ?p ?o .  
    FILTER(isLiteral(?o))  
}
```

Prompt template (type-level):

“Given the following literal predicates observed on dbo:Film, which one most likely contains the movie’s title? Candidates: {rdfs:label, dbp:title, dbo:abstract, …}”

Concrete decision for *type* dbo:Film:

- LLM selects (for dbo:Film): dbp:title (*fallback: rdfs:label if missing*).
- Then for any film entity, use the value of dbp:title in prompts instead of the URI, e.g., “Who directed Inception?”.

4.3 Algorithm 2: Seed Entity Sampling and Subgraph Extraction

Algorithm 2 Entity and Subgraph Extraction For Node Type t

Input: endpoint : SPARQL endpoint, p_l : entity label of t , n : number of entities of t , BZ : entity batch size, h : number of hops for subgraph, shape : Subgraph Shape, d : Direction of predicates

Output: Entity_Subgraph : A list of n valid entity, subgraph pair

```
1:  $\text{entity\_subgraph} \leftarrow \{\}$ 
2: while  $\text{entity\_subgraph.length} < n$  do
3:    $e, e_L \leftarrow \text{getEntity}(\text{endpoint}, p_l, BZ)$ 
4:    $gsize \leftarrow \text{getGraphSize}(\text{endpoint}, e, h)$ 
5:    $\text{preds} \leftarrow \text{countUniquePredicates}(\text{endpoint}, e, h)$ 
6:    $\text{context}_{\text{valid}} \leftarrow \text{validateContext}(gsize, \text{preds})$ 
7:    $\text{subgraph} \leftarrow \text{extractSubgraph}(e, \text{endpoint}, h, \text{shape}, d)$ 
8:    $\text{subgraph}_{\text{filtered}} \leftarrow \text{filterSubgraph}(\text{subgraph}, p_l)$ 
9:    $\text{subgraph}_{\text{valid}} \leftarrow \text{validateSubgraph}(\text{subgraph}_{\text{filtered}})$ 
10:   $\text{entity\_subgraph}[e] \leftarrow \text{subgraph}_{\text{filtered}}$ 
11: end while
12: Return  $\text{entity\_Subgraph}$ 
```

Algorithm 2

Goal: Efficiently sample high-quality seed entities and extract subgraphs.

Key ideas

- **Batch sampling** from SPARQL endpoint (LIMIT/OFFSET) to approximate random selection without heavy latency.
- **Early filtering** on label length, predicate diversity, and subgraph size to avoid low-quality seeds.
- **Predicate blacklist / literal trimming** to remove thumbnails/long abstracts and keep salient facts.
- **Stop early** when per-type quota n_t is met to limit queries and cost.

4.3 Algorithm 2: Steps

- ① **Batch entities of type t :** query SPARQL endpoint with LIMIT BZ OFFSET to fetch candidate entity URIs.
- ② **For each candidate entity e :**
 - Retrieve predicted label e_L (from §4.2).
 - Fetch 1-hop triples (incoming + outgoing).
 - Compute subgraph stats: total triples K_e , unique predicate count P_e .
- ③ **Quality checks:**
 - Label exists and is short enough (avoid extremely long titles).
 - Predicate diversity $P_e \geq k$ (e.g. $k = 3$) to ensure multiple questions.
 - Subgraph token size within model context limit.
- ④ **Accept or reject:** accept $(e, SG(e))$ if checks pass; otherwise discard and continue.
- ⑤ **Repeat batches** until collected n_t seeds for type t or no more candidates.
- ⑥ **Return:** list of accepted seed entity — subgraph pairs for downstream generation.

4.3 Algorithm 2: Complexity & Example “Inception”

Complexity: $O(\tau \cdot (q_{2_cost} + K))$

- τ = number of selected representative types,
- q_{2_cost} = cost per batch SPARQL retrieval from endpoint,
- K = average triples per accepted subgraph.

Example: how “Inception” becomes a seed

- Candidate: <<http://dbpedia.org/resource/Inception>> with label “Inception”.
- Label check: “Inception” length small \Rightarrow **pass**.
- 1-hop extraction yields triples (filtered):
 - (Inception, dbp:title, "Inception")
 - (Inception, dbo:director, Christopher_Nolan)
 - (Inception, dbo:starring, Leonardo_DiCaprio)
 - (Inception, dbo:releaseDate, "2010-07-16")
- $P_e = 4 \geq k (= 3)$ and token estimate < model limit \Rightarrow **accept as seed**.

5.1 Algorithm 3: Subgraph Summarization

Goal: Produce a compact, answer-free subgraph summary for use in **Independent Question Generation**.

Steps:

- ① **Collect predicates** from the seed entity's subgraph.
- ② **Select representative triple** for each predicate (first or sampled).
- ③ **Remove concrete objects (answers)** — retain only entity and predicate, replacing object with placeholder.

Algorithm 3 Summarizing Subgraph Algorithm

Input: SG : Subgraph

Output: SSG : Summarized Subgraph

```
1:  $p_u \leftarrow getUniquePredicates(SG)$ 
2:  $preds\_to\_triples \leftarrow groupTriples(SG, p_u)$ 
3:  $SSG \leftarrow []$ 
4: for every  $\langle p, t_{list} \rangle \in preds\_to\_triples$  do
5:    $t' \leftarrow modifyTriple(e, t_{list}[0])$ 
6:    $SSG \leftarrow ssg \cup t'$ 
7: end for
8: Return  $SSG$ 
```

Algorithm 3

Why summarize?

- Reduce token usage to fit more context in prompts.
- Prevent revealing answers — encourages question generation from structure.
- Avoid embedding literal answers inside SPARQL queries later.

5.2 Independent Question Generation

- Input: serialized subgraph SG_{ser} , entity label e_L , desired number n_q .
- Task: generate $Q' = [Q_1, \dots, Q_{n_q}]$ of independent, answerable questions.

$$Q' = f(\text{SG}_{ser}, e_L, n_q, I_{IQ}) \quad (3)$$

$$(Q', T_q) = f(\text{SG}_{ser}, e_L, n_q, I'_{IQ}) \quad (4)$$

Notes: Eq.(3) = independent-question generator. Eq.(4) = extended output that also returns, for each question, the minimal supporting triple(s) T_q (used later for SPARQL generation).

5.3 SPARQL Query Generation

- Input per question: $(Q_i, T_{q,i})$ — $T_{q,i}$ are the minimal triple(s) supporting Q_i .
- Task: produce SPARQL queries SQ_i that retrieve the answer from the KG.

$$SQ = f(Q', T_q, ISQ) \tag{5}$$

Notes: providing T_q reduces prefix/predicate errors and grounds the LLM into using predicates present in the subgraph.

5.4 Dialogue Generation

- Input: independent-question list Q' and entity label e_L .
- Task: turn Q' into a coherent multi-turn dialogue D (first question standalone; later turns use co-reference / pronouns).

$$D = f(Q', e_L, I_{DG}) \quad (6)$$

Notes: Eq.(6) is the final dialogue transform in the multi-stage pipeline (uses e_L to help pronoun/coref handling).

5.5 Validators: Constraints Summary

- **Question Validator:** each Q_i must be answerable from SG and explicitly reference entity (for Q_1).
- **Query Validator:** SPARQL syntax check + execution match to expected triple(s).
- **Dialogue Validator:** coherence checks (pronoun usage, Q_1 independence, no one-word Qs).
- Retry policy: up to 3 attempts per stage, otherwise drop seed and resample.

6.1 Evaluation Setup

- Benchmarks generated on three real-world KGs:
 - **DBLP, DBpedia, YAGO, MAG**
- Compared systems:
 - **Chatty-Gen**
 - **Maestro** (state-of-the-art KG question benchmark generator)
- Two hardware setups used:
 - High-memory server for SPARQL + closed-source LLMs
 - GPU server for open-source LLMs

6.2 Question Quality: Type Diversity

- Maestro often generates repetitive and boolean-type questions.
- Chatty-Gen generates more balanced, human-like question types.
- This improves dialogue naturalness and domain coverage.






Figure: Comparison of question type diversity. (Figure 4 in paper)

6.3 Node Type Coverage in KGs

- Chatty-Gen's entity sampling preserves the original node-type distribution of KG.
- Maestro's selection misses important entity types (e.g., Person in DBpedia).






Figure: Seed node type distribution comparison. (Figure 5 in paper)

6.4 End-to-End Time Efficiency (Table 3)

- Maestro requires full KG preprocessing, scaling poorly on large KGs.
- Chatty-Gen retrieves subgraphs on demand, enabling significant speedups.

| KG | Maestro (hrs) | Chatty-Gen (hrs) |
|---------|---------------|------------------|
| DBpedia | 30.77 | 0.17 |
| YAGO | 5.20 | 0.10 |
| MAG | 5.38 | 0.12 |
| DBLP | 0.12 | 0.12 |

Table: End-to-end generation time (Table 3 in paper).

6.5 Consistent Performance Across LLMs

- Multi-stage pipeline significantly reduces hallucinations and SPARQL errors.
- Open-source setups (e.g., LLaMA + CodeLLaMA) reach performance **comparable to GPT-4o**.

| | Approach | LLM | YAGO | | | | | | DBLP | | | | | | # Tokens M (\$) | | |
|------------------------------|-----------------|-----|----------------|------------|------------|----------|------------|-----------|------------|----------------|------------|------------|----------|------------|-----------------|------------|------------|
| | | | Success rate % | Dialogue-S | Question-E | SPARQL-E | Dialogue-E | Parsing-E | Time(mins) | Success rate % | Dialogue-S | Question-E | SPARQL-E | Dialogue-E | Parsing-E | | |
| Our Multi-stages (3 prompts) | GPT-4o | 100 | 20 | 0 | 0 | 0 | 0 | 0 | 8 | 0.04(0.32) | 95 | 20 | 0 | 0 | 1 | 0.05(0.38) | |
| | GPT-4 | 95 | 20 | 0 | 0 | 1 | 0 | 13 | 0.05(0.71) | 100 | 20 | 0 | 0 | 0 | 11 | 0.05(0.78) | |
| | GPT-3.5 | 91 | 20 | 1 | 0 | 1 | 0 | 6 | 0.05(0.04) | 95 | 20 | 1 | 0 | 0 | 0 | 7 | 0.06(0.05) |
| | Gemini-1-pro | 71 | 20 | 0 | 2 | 0 | 6 | 6 | 0.06(0.01) | 67 | 20 | 5 | 0 | 0 | 5 | 5.2 | 0.09(0.02) |
| | Gemini-1.5-pro | 22 | 20 | 0 | 1 | 0 | 71 | 20.5 | 0.18(0.35) | 41 | 20 | 0 | 0 | 1 | 28 | 14.5 | 0.12(0.24) |
| | LLAMA-3-8b | 13 | 20 | 2 | 124 | 9 | 0 | 84 | 0.36 | 20 | 20 | 2 | 47 | 23 | 8 | 71 | 0.33 |
| | LLAMA-3-8b-inst | 41 | 20 | 0 | 16 | 0 | 13 | 37 | 0.13 | 14 | 20 | 0 | 106 | 0 | 16 | 89 | 0.39 |
| | LLAMA-2-13b | 5 | 20 | 5 | 325 | 10 | 1 | 270 | 0.79 | 1 | 5 | 32 | 426 | 7 | 22 | 392 | 1.30 |
| | CodeLLAMA-7b | 7 | 20 | 3 | 162 | 74 | 16 | 234 | 0.92 | 1 | 5 | 144 | 156 | 44 | 145 | 356 | 0.16 |
| | CodeLLAMA-13b | 83 | 20 | 0 | 0 | 4 | 0 | 31 | 0.08 | 63 | 20 | 3 | 5 | 2 | 1 | 37 | 0.01 |
| single prompt | Mistral-7b-v0.1 | 17 | 20 | 2 | 88 | 3 | 0 | 70 | 0.25 | 3 | 15 | 5 | 323 | 5 | 14 | 189 | 0.90 |
| | Multi-LLM-1 | 100 | 20 | 0 | 0 | 0 | 0 | 18 | 0.06 | 83 | 20 | 0 | 4 | 0 | 0 | 20 | 0.07 |
| | Multi-LLM-2 | 91 | 20 | 0 | 0 | 0 | 2 | 18 | 0.06(0.28) | 71 | 20 | 2 | 4 | 0 | 2 | 23 | 0.10(0.82) |
| | GPT-4o | 38 | 20 | 11 | 0 | 5 | 15 | 9 | 0.03(0.32) | 8 | 20 | 52 | 0 | 8 | 144 | 59 | 0.15(1.30) |
| | GPT-4 | 40 | 20 | 5 | 1 | 1 | 23 | 17 | 0.04(0.73) | 21 | 20 | 5 | 1 | 2 | 66 | 46 | 0.09(1.86) |
| | GPT-3.5 | 5 | 19 | 2 | 0 | 8 | 242 | 53 | 0.36(0.32) | 8 | 20 | 25 | 0 | 5 | 119 | 50 | 0.25(0.24) |
| | Gemini-1-pro | 56 | 20 | 0 | 10 | 4 | 1 | 3 | 0.03(0.01) | 10 | 20 | 130 | 0 | 12 | 6 | 21 | 0.30(0.07) |
| | Gemini-1.5-pro | 14 | 20 | 59 | 4 | 0 | 5 | 16 | 0.12(0.27) | 6 | 20 | 28 | 0 | 3 | 10 | 51.5 | 0.43(0.9) |
| | LLAMA-3-8b | 0 | 0 | 44 | 107 | 158 | 92 | 251 | 1 | 0 | 2 | 170 | 23 | 55 | 452 | 286 | 1.23 |
| | LLAMA-3-8b-inst | 9 | 20 | 10 | 13 | 59 | 43 | 72 | 0.26 | 7 | 20 | 51 | 6 | 44 | 71 | 90 | 0.38 |
| | LLAMA-2-13b | 0 | 0 | 40 | 216 | 87 | 401 | 403 | 1 | 0 | 0 | 53 | 39 | 18 | 827 | 479 | 1.22 |
| | CodeLLAMA-7b | 0 | 0 | 72 | 87 | 196 | 344 | 314 | 1.07 | 0 | 1 | 203 | 45 | 30 | 604 | 313 | 1.25 |
| | CodeLLAMA-13b | 1 | 11 | 78 | 29 | 579 | 168 | 443 | 1.07 | 1 | 5 | 188 | 11 | 91 | 677 | 459 | 1.26 |
| | Mistral-7b-v0.1 | 0 | 0 | 53 | 93 | 420 | 223 | 261 | 0.98 | 0 | 1 | 208 | 34 | 118 | 528 | 257 | 1.20 |

Figure: LLM performance comparison (Table 4 in the paper).

6.6 Summarized vs Full Graphs

- Compare using **Summarized subgraphs** vs **Full subgraphs**.
- Summarization reduces prompt size → fewer tokens → lower cost.
- At the same time, it **improves SPARQL correctness**.




Figure: Token usage and SPARQL correctness comparison (Figure 6 in the paper).

7.1 Conclusion

- Chatty-Gen presents a cost-effective pipeline for generating **multi-turn dialogue benchmarks** directly from KGs.
- The system **avoids full KG preprocessing** and instead performs **on-demand subgraph retrieval**, significantly reducing generation time.
- By using a **multi-stage LLM prompting strategy** and optional **subgraph summarization**, Chatty-Gen enables **consistent performance across both proprietary and open-source LLMs**.
- Evaluations show that Chatty-Gen generates dialogues that are **more diverse, entity-grounded, and semantically aligned** with the original KG compared to prior systems (e.g., Maestro).

7.2 Limitations

1. Limited Question Complexity

- Chatty-Gen relies on 1-hop subgraphs and predicate summarization.
- Generated questions are mainly factual (What/Who/When) — lack multi-hop reasoning or comparative logic.
- Future work: adaptive multi-hop retrieval and reasoning-path prompting.

2. Error Propagation in Multi-Stage Pipeline

- Each stage (entity representation → question → SPARQL → dialogue) depends on previous LLM outputs.
- A single semantic or predicate error propagates through later stages, degrading overall quality.
- Validators mitigate syntax-level errors but cannot fully eliminate semantic drift.